



CLOUD  
USER **24**  
MEETING



## Advanced Kubernetes Day 1: Workflows on K8s

**Sebastian Beyvers** | Bioinformatics & Systems Biology | JLU Giessen

[sebastian.beyvers@cb.jlug.de](mailto:sebastian.beyvers@cb.jlug.de)

04.12.2024



# Cloud Native Workflow Engines

- Kubernetes native workflow engine
- Uses custom resources and the “operator pattern”
- Containerized steps
- DAG definition via yamls
- Built-in scripting support
- Can create other K8s resources
- Data science focused Python library (hera)

## argoproj/argo-workflows

Workflow Engine for Kubernetes



856

Contributors

356

Used by

1k

Discussions

15k

Stars

3k

Forks



# Simple example

Containers

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: hello-world- # Name of this Workflow
spec:
  entrypoint: hello-world    # Defines "hello-world" as the "main" template
  templates:
    - name: hello-world      # Defines the "hello-world" template
      container:
        image: busybox
        command: [echo]
        args: ["hello world"] # This template runs "echo" in the "busybox" image w
```

Scripts

```
- name: gen-random-int
  script:
    image: python:alpine3.6
    command: [python]
    source: |
      import random
      i = random.randint(1, 100)
      print(i)
```

Resources

```
- name: k8s-owner-reference
  resource:
    action: create
    manifest: |
      apiVersion: v1
      kind: ConfigMap
      metadata:
        generateName: owned-eg-
      data:
        some: value
```

# Templates

## Step based workflows

```
- name: hello-hello-hello
  steps:
  - - name: step1
      template: prepare-data
  - - name: step2a
      template: run-data-first-half
    - name: step2b
      template: run-data-second-half
```

## DAG based workflow

```
- name: diamond
  dag:
    tasks:
    - name: A
      template: echo
    - name: B
      dependencies: [A]
      template: echo
    - name: C
      dependencies: [A]
      template: echo
    - name: D
      dependencies: [B, C]
      template: echo
```

# Nextflow on Kubernetes: Old way

- Shared volume needed for nextflow work directory
- All steps must be a container
- Easiest to spawn from inside the cluster

```
process {  
    executor = 'k8s'  
}  
  
k8s {  
    storageClaimName = 'vol-claim'  
    storageMountPath = '/mount/path'  
    storageSubPath = '/my-data'  
}
```

Nextflow config

# Nextflow on Kubernetes

- New:
  - Uses a fuse s3 adapter
    - Needs *privileged* permissions / or a custom controller
  - requires S3 configuration
  - Uses “weave” containers that can be automatically built based on conda environments

```
wave {  
    enabled = true  
}  
  
fusion {  
    enabled = true  
}  
  
process {  
    executor = 'k8s'  
}  
  
k8s {  
    context = '<YOUR K8S CONFIGURATION CONTEXT>'  
    namespace = '<YOUR K8S NAMESPACE>'  
    serviceAccount = '<YOUR K8S SERVICE ACCOUNT>'  
}
```

**Hands on !**

# Snakemake on Kubernetes

- Needs an external storage setup (like S3)
- Preferably needs pre-configured workflows from git
- All intermediate results will be put into S3 storage
- Built-in conda support to create environments on the fly
- Unfortunately: Very sparsely documented





**Hands on !**

# A full service demonstration

- Kubernetes native workflow engine
- Uses custom resources and the “operator pattern”
- Containerized steps
- DAG definition via yamls
- Built-in scripting support
- Can create other K8s resources
- Data science focused Python library (hera)

## argoproj/argo-workflows

Workflow Engine for Kubernetes



856

Contributors

356

Used by

1k

Discussions

15k

Stars

3k

Forks



**Questions ?**